

Docker

an introduction

Markus van Dijk



Process isolation

Why

- Testing: reproducible environment
- Multiple instances
- Library versions
- Vulnerable SW

Process isolation

History

- Unix `chroot` (1982)
- BSD `jail` (2000)
- Linux containers LXC (2008)
- Docker (2013)
- Open Container Initiative (2015)

But we have VMs

Containers

- Shared OS
- FS copy-on-write
- Lightweight start/stop
- *microservice*
- One container, one job
- External plumbing
- Prepare to fail

VMs

- Independent OSes
- Duplicated FS
- Full boot/shutdown
- *LAMP* etc
- Many jobs
- Configure firewalls etc
- Robust

Typical use

- Develop and deploy in same environment
Grid jobs: test on laptop!
- Task-driven service activation
Front-end starts container for each job
- Easy migration iff prepare to fail
Kill containers, restart jobs in new location
- Connected containers, need not run on same host
DB servers, web front, job processors

Isolation

- File system
 - Processes
 - Network
 - User IDs (coming soon)
-
- Some vulnerabilities
Run in VM for complete protection for now

Docker toolbox

- **docker**
Defines a single container and manages its lifecycle
View logs, execute a process in running container
- **docker-machine**
Transparently runs docker commands on remote host
- **docker-compose**
Defines a set of connected containers, manages dependencies in lifecycle
- Various tools for orchestration
fleet, swarm, mesos, ...

Example: Cloud9

<https://c9.io>

- Cloud-based development workspace
- 1 docker workspace per project
- Ssh root access to your workspace
- Web browser IDE
- Container down after idle period
- Container up in seconds
- Project footprint: only your files

